



A Indexing and reduction overhead: executed instructions (same device, profiled input)




† GH200 Triton uses a different autotuned configuration (confounder)


B GH200 matmul: shared-memory stores per WGMMMA




B200 matmul: TMA operand load bytes




MI300X matmul (diagnostic): descriptor vs. pointer loads




C 1d_conv FP16: load traffic



1d_conv FP16: achieved and theoretical occupancy



MI300X vector_add (diagnostic): load modifier × cache flush



issue % 58 78 4 61 66 22